

# Complete Guide to Recurrent Neural Networks

## RNN, LSTM, GRU, and Advanced LSTM Architectures

---

### Table of Contents

1. [Introduction to Sequential Data](#)
  2. [Recurrent Neural Networks \(RNN\)](#)
  3. [Long Short-Term Memory \(LSTM\)](#)
  4. [Gated Recurrent Unit \(GRU\)](#)
  5. [Advanced LSTM Architectures](#)
  6. [Detailed Comparisons](#)
  7. [Implementation Examples](#)
  8. [Use Case Scenarios](#)
  9. [Performance Analysis](#)
  10. [Best Practices](#)
- 

## 1. Introduction to Sequential Data {#introduction}

### What is Sequential Data?

Sequential data is information where the order of elements matters. Examples include:

- **Text:** "I love this movie" vs "This movie I love" (different meanings)
- **Time Series:** Stock prices, weather data, sensor readings
- **Speech:** Audio signals where timing is crucial
- **DNA Sequences:** Order of nucleotides determines genetic information

### Why Traditional Neural Networks Fail?

Standard feedforward neural networks treat each input independently:

Input: [word1, word2, word3]

Processing:  $f(\text{word1})$ ,  $f(\text{word2})$ ,  $f(\text{word3})$  - No connection between words

### Problems:

- No memory of previous inputs
- Fixed input size requirement
- Cannot capture temporal dependencies
- Loses context information

## The Need for Recurrent Networks

Recurrent networks process sequences by maintaining **hidden states** that carry information from previous time steps:

$$\begin{array}{ccccccc}
 h_0 & \rightarrow & h_1 & \rightarrow & h_2 & \rightarrow & h_3 \rightarrow \dots \\
 \uparrow & & \uparrow & & \uparrow & & \uparrow \\
 x_1 & & x_2 & & x_3 & & x_4
 \end{array}$$

## 2. Recurrent Neural Networks (RNN) {#rnn}

### 2.1 Basic RNN Architecture

#### Mathematical Foundation

$$\begin{aligned}
 h_t &= \tanh(W_{hh} \times h_{t-1} + W_{xh} \times x_t + b_h) \\
 y_t &= W_{hy} \times h_t + b_y
 \end{aligned}$$

#### Where:

- $h_t$ : Hidden state at time t
- $x_t$ : Input at time t
- $y_t$ : Output at time t
- $W_{hh}$ : Hidden-to-hidden weight matrix
- $W_{xh}$ : Input-to-hidden weight matrix
- $W_{hy}$ : Hidden-to-output weight matrix
- $b_h, b_y$ : Bias vectors

#### Visual Representation

$y_1$	$y_2$	$y_3$
$\uparrow$	$\uparrow$	$\uparrow$
$[h_1] \rightarrow [h_2] \rightarrow [h_3]$		
$\uparrow$	$\uparrow$	$\uparrow$
$x_1$	$x_2$	$x_3$

## 2.2 RNN Processing Flow

### Step-by-Step Example: Sentiment Analysis

**Input Sentence:** "This movie is great"

**Step 1:** Initialize  $h_0 = [0, 0, 0, 0]$  (4-dimensional hidden state)

**Step 2:** Process "This"

```
x1 = embedding("This") = [0.2, -0.1, 0.5, 0.3]
h1 = tanh(W_hh × h0 + W_xh × x1 + b_h)
h1 = [0.1, 0.3, -0.2, 0.4]
```

**Step 3:** Process "movie"

```
x2 = embedding("movie") = [0.4, 0.1, -0.3, 0.2]
h2 = tanh(W_hh × h1 + W_xh × x2 + b_h)
h2 = [0.3, 0.2, 0.1, 0.5]
```

**Continue for remaining words...**

## 2.3 Strengths of RNN

1. **Variable Length Input:** Can handle sequences of any length
2. **Parameter Sharing:** Same weights used across all time steps
3. **Memory:** Maintains information from previous inputs
4. **Contextual Understanding:** Can capture dependencies between elements

## 2.4 Critical Limitations

### Vanishing Gradient Problem

**Mathematical Explanation:** During backpropagation through time, gradients are computed as:

$$\partial L / \partial h_t = \partial L / \partial h_{t+1} \times \partial h_{t+1} / \partial h_t$$

Since  $\partial h_{t+1} / \partial h_t$  involves the tanh derivative (max value = 1), gradients shrink exponentially:

$$\partial L / \partial h_1 = \partial L / \partial h_T \times (\partial h_T / \partial h_{T-1}) \times \dots \times (\partial h_2 / \partial h_1)$$

**Consequence:** Earlier time steps receive negligible gradient updates.

## Practical Example of Vanishing Gradients

Sentence: "The cat, which was sitting on the mat that my grandmother bought years ago, was sleeping."

Problem: By the time RNN processes "was sleeping", it has forgotten "cat" due to vanishing gradients.

Result: Poor understanding of long-range dependencies.

## Short-Term Memory Problem

RNNs struggle with information that occurred many steps ago:

Text: "I grew up in France... [100 words later] ...so I speak French fluently."

RNN often fails to connect "France" with "French" due to the gap.

## 2.5 RNN Variants

### Bidirectional RNN

Processes sequences in both directions:

Forward:  $h_1 \rightarrow h_2 \rightarrow h_3 \rightarrow h_4$

Backward:  $h_4 \leftarrow h_3 \leftarrow h_2 \leftarrow h_1$

**Advantage:** Access to both past and future context **Use Case:** When entire sequence is available (not real-time)

### Deep RNN

Multiple RNN layers stacked:

Layer 2:  $h_1^2 \rightarrow h_2^2 \rightarrow h_3^2$

↑    ↑    ↑

Layer 1:  $h_1^1 \rightarrow h_2^1 \rightarrow h_3^1$

↑    ↑    ↑

Input:  $x_1$     $x_2$     $x_3$

### 3. Long Short-Term Memory (LSTM) {#lstm}

#### 3.1 The LSTM Solution

LSTM addresses RNN's vanishing gradient problem through a sophisticated **gating mechanism** that controls information flow.

##### Key Innovation: Cell State

$C_t$ : Cell state (long-term memory)

$h_t$ : Hidden state (short-term memory)

#### 3.2 LSTM Architecture Components

##### 3.2.1 Forget Gate

**Purpose:** Decide what information to discard from cell state

$$f_t = \sigma(W_f \times [h_{t-1}, x_t] + b_f)$$

##### Example in Sentiment Analysis:

Text: "The movie started boring but became exciting"

When processing "exciting", forget gate might decide to forget the negative sentiment from "boring"

##### 3.2.2 Input Gate

**Purpose:** Decide what new information to store

$$i_t = \sigma(W_i \times [h_{t-1}, x_t] + b_i)$$

$$\tilde{C}_t = \tanh(W_C \times [h_{t-1}, x_t] + b_C)$$

##### Example:

When processing "exciting", input gate decides to store this positive sentiment information

### 3.2.3 Cell State Update

**Purpose:** Update long-term memory

$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$

**Intuition:**

- Forget irrelevant old information ( $f_t * C_{t-1}$ )
- Add relevant new information ( $i_t * \tilde{C}_t$ )

### 3.2.4 Output Gate

**Purpose:** Control what parts of cell state to output

$$o_t = \sigma(W_o \times [h_{t-1}, x_t] + b_o)$$
$$h_t = o_t * \tanh(C_t)$$

## 3.3 Complete LSTM Forward Pass

### Mathematical Formulation

$$f_t = \sigma(W_f \times [h_{t-1}, x_t] + b_f) \quad \# \text{ Forget gate}$$
$$i_t = \sigma(W_i \times [h_{t-1}, x_t] + b_i) \quad \# \text{ Input gate}$$
$$\tilde{C}_t = \tanh(W_C \times [h_{t-1}, x_t] + b_C) \quad \# \text{ Candidate values}$$
$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t \quad \# \text{ Cell state}$$
$$o_t = \sigma(W_o \times [h_{t-1}, x_t] + b_o) \quad \# \text{ Output gate}$$
$$h_t = o_t * \tanh(C_t) \quad \# \text{ Hidden state}$$

### Visual Flow Diagram



RNN:  $\partial h_t / \partial h_{t-1} = W \times \tanh'(\dots)$  (multiplicative, diminishing)

LSTM:  $\partial C_t / \partial C_{t-1} = f_t$  (additive path available)

## 3.5 LSTM Variants

### 3.5.1 Peephole Connections

Allow gates to look at cell state:

$$f_t = \sigma(W_f \times [C_{t-1}, h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \times [C_{t-1}, h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma(W_o \times [C_t, h_{t-1}, x_t] + b_o)$$

### 3.5.2 Coupled Forget and Input Gates

$$f_t = \sigma(W_f \times [h_{t-1}, x_t] + b_f)$$

$$i_t = 1 - f_t$$

## 4. Gated Recurrent Unit (GRU) {#gru}

### 4.1 GRU Motivation

GRU simplifies LSTM by:

- Combining cell and hidden states
- Using only 2 gates instead of 3
- Reducing computational complexity

### 4.2 GRU Architecture

#### Mathematical Formulation

$$z_t = \sigma(W_z \times [h_{t-1}, x_t]) \quad \# \text{ Update gate}$$

$$r_t = \sigma(W_r \times [h_{t-1}, x_t]) \quad \# \text{ Reset gate}$$

$$\tilde{h}_t = \tanh(W \times [r_t * h_{t-1}, x_t]) \quad \# \text{ Candidate hidden state}$$

$$h_t = (1 - z_t) * h_{t-1} + z_t * \tilde{h}_t \quad \# \text{ Final hidden state}$$

#### Gate Functions Explained

**Update Gate ( $z_t$ ):** Controls how much past information to keep



- $z_t \approx 0$ : Keep old information ( $h_{t-1}$ )
- $z_t \approx 1$ : Use new information ( $\tilde{h}_t$ )

**Reset Gate ( $r_t$ ):** Controls how much past information to forget when computing candidate

- $r_t \approx 0$ : Ignore past hidden state
- $r_t \approx 1$ : Use full past hidden state

## 4.3 GRU vs LSTM Comparison

### Structural Differences

LSTM: Separate cell state ( $C_t$ ) and hidden state ( $h_t$ )

GRU: Combined state ( $h_t$  only)

LSTM: 3 gates (forget, input, output)

GRU: 2 gates (update, reset)

### Parameter Count

LSTM:  $4 \times (\text{input\_size} + \text{hidden\_size} + 1) \times \text{hidden\_size}$  parameters

GRU:  $3 \times (\text{input\_size} + \text{hidden\_size} + 1) \times \text{hidden\_size}$  parameters

Example with  $\text{input\_size}=100$ ,  $\text{hidden\_size}=128$ :

LSTM:  $4 \times (100 + 128 + 1) \times 128 = 117,504$  parameters

GRU:  $3 \times (100 + 128 + 1) \times 128 = 88,128$  parameters

## 4.4 GRU Processing Example

### Sentiment Analysis: "This movie is not good"

#### Step 1: Process "This"

$z_1 = \sigma([0.6, 0.4, 0.2, 0.8])$  # Update gate

$r_1 = \sigma([0.3, 0.7, 0.5, 0.9])$  # Reset gate

$\tilde{h}_1 = \tanh([0.1, 0.4, -0.2, 0.6])$  # Candidate

$h_1 = (1-z_1) * h_0 + z_1 * \tilde{h}_1$  # Combined state

#### Step 2: Process "movie"

Context: Now knows about "This"  
Updates: Incorporates "movie" information

### Step 3: Process "is"

Function: Minimal content, gates likely reduce influence

### Step 4: Process "not"

Critical: Reset gate might prepare to change sentiment direction

### Step 5: Process "good"

Integration: Update gate balances "good" with "not" context  
Result: Overall negative sentiment despite "good"

---

## 5. Advanced LSTM Architectures {#advanced-lstm}

### 5.1 Stacked LSTM

#### Architecture

Layer 3:  $LSTM_3 \rightarrow h_{3,t}$   
↑  
Layer 2:  $LSTM_2 \rightarrow h_{2,t}$   
↑  
Layer 1:  $LSTM_1 \rightarrow h_{1,t}$   
↑  
Input:  $x_t$

#### Benefits

- **Hierarchical Representation:** Each layer captures different abstraction levels
- **Increased Capacity:** More parameters for complex pattern recognition
- **Feature Extraction:** Lower layers extract basic features, higher layers combine them

#### Layer Interpretation Example (Text Processing)

Layer 1: Word-level features (POS tags, word types)  
Layer 2: Phrase-level patterns (noun phrases, verb phrases)  
Layer 3: Sentence-level semantics (sentiment, intent)

## 5.2 Bidirectional LSTM

### Architecture

Forward LSTM:  $h_1^f \rightarrow h_2^f \rightarrow h_3^f \rightarrow h_4^f$   
Backward LSTM:  $h_4^b \leftarrow h_3^b \leftarrow h_2^b \leftarrow h_1^b$   
Combined:  $[h_1^f; h_1^b] [h_2^f; h_2^b] [h_3^f; h_3^b] [h_4^f; h_4^b]$

### Mathematical Formulation

$h_t^f = \text{LSTM\_forward}(x_t, h_{t-1}^f)$  # Forward pass  
 $h_t^b = \text{LSTM\_backward}(x_t, h_{t+1}^b)$  # Backward pass  
 $h_t = [h_t^f; h_t^b]$  # Concatenation

### Use Case Example: Named Entity Recognition

Sentence: "Barack Obama was born in Hawaii"

Forward context at "Obama": ["Barack"]  
Backward context at "Obama": ["was", "born", "in", "Hawaii"]  
Combined understanding: "Barack Obama" is likely a person name

## 5.3 LSTM with Attention Mechanism

### Attention Concept

Instead of using only the final hidden state, attention allows the model to focus on relevant parts of the entire sequence.

### Architecture

LSTM Hidden States:  $h_1, h_2, h_3, h_4$   
Attention Weights:  $\alpha_1, \alpha_2, \alpha_3, \alpha_4$  (sum to 1)  
Context Vector:  $c = \sum(\alpha_i \times h_i)$

### Mathematical Formulation

```
e_i = score(h_i, h_final)      # Attention scores
α_i = softmax(e_i)             # Normalized weights
c = Σ(α_i × h_i)               # Weighted sum
```

## Example: Machine Translation

English: "The cat sits on the mat"  
Hidden states:  $[h_1, h_2, h_3, h_4, h_5, h_6]$

When translating to "El gato":  
- High attention on  $h_1$  ("The") and  $h_2$  ("cat")  
- Low attention on other positions

## 5.4 LSTM with Dropout

### Dropout Variants

#### 1. Standard Dropout: Applied to non-recurrent connections

```
python
x_t = dropout(x_t, rate=0.2)
h_t = LSTM(x_t, h_{t-1})
```

#### 2. Recurrent Dropout: Applied to recurrent connections

```
python
h_{t-1} = recurrent_dropout(h_{t-1}, rate=0.2)
h_t = LSTM(x_t, h_{t-1})
```

#### 3. Zoneout: Randomly keeps some hidden units unchanged

```
python
mask = random_binary_mask(rate=0.1)
h_t = mask * h_{t-1} + (1 - mask) * h_t_new
```

---

## 6. Detailed Comparisons {#comparisons}

### 6.1 Performance Comparison

## Computational Complexity

Model	Parameters	Training Time	Memory Usage	Inference Speed
RNN	$O(W^2)$	Fast	Low	Fastest
LSTM	$O(4W^2)$	Slow	High	Slow
GRU	$O(3W^2)$	Medium	Medium	Medium
Advanced LSTM	$O(8W^2+)$	Slowest	Highest	Slowest

Where  $W = \text{hidden\_size} + \text{input\_size}$

## Memory Requirements (Exact Calculations)

Given:  $\text{input\_size} = 100$ ,  $\text{hidden\_size} = 128$ ,  $\text{sequence\_length} = 200$

RNN Memory:

- Weights:  $(100 + 128) \times 128 = 29,184$  parameters
- Hidden states:  $200 \times 128 = 25,600$  values
- Total: ~54K values

LSTM Memory:

- Weights:  $4 \times (100 + 128 + 1) \times 128 = 117,504$  parameters
- Hidden + Cell states:  $200 \times 128 \times 2 = 51,200$  values
- Gate activations:  $200 \times 128 \times 4 = 102,400$  values
- Total: ~271K values

GRU Memory:

- Weights:  $3 \times (100 + 128 + 1) \times 128 = 88,128$  parameters
- Hidden states:  $200 \times 128 = 25,600$  values
- Gate activations:  $200 \times 128 \times 3 = 76,800$  values
- Total: ~190K values

## 6.2 Accuracy Comparison by Task

### Task 1: Sentiment Analysis (IMDB)

Dataset: 50,000 movie reviews

Metric: Accuracy

Results:

- Simple RNN:  $0.831 \pm 0.012$
- LSTM:  $0.871 \pm 0.008$
- GRU:  $0.867 \pm 0.009$
- Bidirectional LSTM:  $0.884 \pm 0.006$
- Stacked LSTM:  $0.892 \pm 0.007$

## Task 2: Language Modeling (Penn Treebank)

Dataset: Penn Treebank corpus

Metric: Perplexity (lower is better)

Results:

- Simple RNN: 165.2
- LSTM: 78.4
- GRU: 81.9
- Advanced LSTM: 68.7

## Task 3: Machine Translation (WMT14)

Dataset: English-German translation

Metric: BLEU Score (higher is better)

Results:

- LSTM: 24.8
- Bidirectional LSTM: 26.3
- LSTM + Attention: 28.1
- Stacked LSTM + Attention: 29.7

## 6.3 Training Characteristics

### Convergence Speed

Epochs to 90% of final performance:

Simple RNN: 15-20 epochs (but lower final performance)

LSTM: 25-30 epochs

GRU: 20-25 epochs

Advanced LSTM: 35-45 epochs (but higher final performance)

Gradient Stability

Gradient Norm Statistics over Training:

RNN: High variance, frequent exploding/vanishing

LSTM: Stable, consistent gradients

GRU: Stable, slightly more variance than LSTM

Advanced LSTM: Very stable, slower but consistent

6.4 Sequence Length Handling

Performance vs Sequence Length

Sequence Length: 50 100 200 500 1000  
RNN Accuracy: 0.85 0.78 0.65 0.45 0.30  
LSTM Accuracy: 0.87 0.86 0.84 0.79 0.72  
GRU Accuracy: 0.86 0.85 0.83 0.78 0.70  
Adv LSTM Acc: 0.89 0.88 0.87 0.83 0.78

6.5 Architecture Decision Matrix

Criteria	RNN	LSTM	GRU	Advanced LSTM
Short sequences (<50)	✔ Good	✔ Good	✔ Good	⚠ Overkill
Medium sequences (50-200)	⚠ Limited	✔ Excellent	✔ Excellent	✔ Best
Long sequences (>200)	✖ Poor	✔ Good	✔ Good	✔ Excellent
Limited compute	✔ Best	⚠ Moderate	✔ Good	✖ Poor
High accuracy needed	✖ Limited	✔ Good	✔ Good	✔ Best
Real-time inference	✔ Fastest	⚠ Slow	✔ Moderate	✖ Slowest
Small datasets	✔ Good	⚠ Overfitting	✔ Good	✖ Overfitting
Large datasets	⚠ Limited	✔ Excellent	✔ Excellent	✔ Best

## 7. Implementation Examples {#examples}

### 7.1 Simple RNN Implementation

#### PyTorch Implementation

```
python

import torch
import torch.nn as nn

class SimpleRNN(nn.Module):
    def __init__(self, input_size, hidden_size, output_size, num_layers=1):
        super(SimpleRNN, self).__init__()
        self.hidden_size = hidden_size
        self.num_layers = num_layers

        # RNN layer
        self.rnn = nn.RNN(input_size, hidden_size, num_layers,
                           batch_first=True)

        # Output layer
        self.fc = nn.Linear(hidden_size, output_size)

    def forward(self, x):
        # Initialize hidden state
        h0 = torch.zeros(self.num_layers, x.size(0), self.hidden_size)

        # Forward propagate RNN
        out, _ = self.rnn(x, h0)

        # Use the last output for classification
        out = self.fc(out[:, -1, :])
        return out

# Usage example
model = SimpleRNN(input_size=100, hidden_size=128, output_size=2)
input_tensor = torch.randn(32, 50, 100) # (batch, seq_len, features)
output = model(input_tensor)
print(f"Output shape: {output.shape}") # [32, 2]
```

#### TensorFlow/Keras Implementation

```
python
```



```
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import SimpleRNN, Dense, Dropout

def create_simple_rnn(input_shape, hidden_units, output_units):
    model = Sequential([
        SimpleRNN(hidden_units,
                   return_sequences=False,
                   dropout=0.2,
                   recurrent_dropout=0.2,
                   input_shape=input_shape),
        Dense(64, activation='relu'),
        Dropout(0.5),
        Dense(output_units, activation='sigmoid')
    ])

    model.compile(optimizer='adam',
                  loss='binary_crossentropy',
                  metrics=['accuracy'])
    return model

# Usage
model = create_simple_rnn((200, 100), 128, 1)
model.summary()
```

## 7.2 LSTM Implementation

### Detailed LSTM with Custom Forward Pass

```
python
```

```

import torch
import torch.nn as nn
import torch.nn.functional as F

class CustomLSTM(nn.Module):
    def __init__(self, input_size, hidden_size):
        super(CustomLSTM, self).__init__()
        self.input_size = input_size
        self.hidden_size = hidden_size

        # Gate parameters
        self.W_f = nn.Linear(input_size + hidden_size, hidden_size) # Forget gate
        self.W_i = nn.Linear(input_size + hidden_size, hidden_size) # Input gate
        self.W_C = nn.Linear(input_size + hidden_size, hidden_size) # Candidate
        self.W_o = nn.Linear(input_size + hidden_size, hidden_size) # Output gate

    def forward(self, x, hidden=None):
        batch_size, seq_len, _ = x.size()

        if hidden is None:
            h_t = torch.zeros(batch_size, self.hidden_size)
            C_t = torch.zeros(batch_size, self.hidden_size)
        else:
            h_t, C_t = hidden

        outputs = []

        for t in range(seq_len):
            x_t = x[:, t, :]

            # Concatenate input and hidden state
            combined = torch.cat([x_t, h_t], dim=1)

            # Gate computations
            f_t = torch.sigmoid(self.W_f(combined)) # Forget gate
            i_t = torch.sigmoid(self.W_i(combined)) # Input gate
            C_tilde_t = torch.tanh(self.W_C(combined)) # Candidate values
            o_t = torch.sigmoid(self.W_o(combined)) # Output gate

            # Update cell state
            C_t = f_t * C_t + i_t * C_tilde_t

            # Update hidden state

```

```
h_t = o_t * torch.tanh(C_t)
```

```
outputs.append(h_t.unsqueeze(1))
```

```
outputs = torch.cat(outputs, dim=1)
```

```
return outputs, (h_t, C_t)
```

*# Advanced LSTM for Sentiment Analysis*

```
class SentimentLSTM(nn.Module):
```

```
    def __init__(self, vocab_size, embedding_dim, hidden_dim, num_layers, dropout=0.5):
```

```
        super(SentimentLSTM, self).__init__()
```

*# Embedding layer*

```
self.embedding = nn.Embedding(vocab_size, embedding_dim)
```

*# LSTM layers*

```
self.lstm = nn.LSTM(embedding_dim, hidden_dim, num_layers,  
                    dropout=dropout, batch_first=True,  
                    bidirectional=True)
```

*# Attention mechanism*

```
self.attention = nn.Linear(hidden_dim * 2, 1)
```

*# Classification layers*

```
self.fc1 = nn.Linear(hidden_dim * 2, hidden_dim)
```

```
self.dropout = nn.Dropout(dropout)
```

```
self.fc2 = nn.Linear(hidden_dim, 1)
```

```
def attention_mechanism(self, lstm_out):
```

*# lstm\_out: (batch\_size, seq\_len, hidden\_dim \* 2)*

```
attention_weights = F.softmax(self.attention(lstm_out).squeeze(-1), dim=1)
```

*# attention\_weights: (batch\_size, seq\_len)*

*# Apply attention weights*

```
context_vector = torch.sum(lstm_out * attention_weights.unsqueeze(-1), dim=1)
```

```
return context_vector, attention_weights
```

```
def forward(self, x):
```

*# Embedding*

```
embedded = self.embedding(x)
```

*# LSTM*

```
lstm_out, _ = self.lstm(embedded)
```

*# Attention*

```
context_vector, attention_weights = self.attention_mechanism(lstm_out)
```

*# Classification*

```
out = F.relu(self.fc1(context_vector))
```

```
out = self.dropout(out)
```

```
out = torch.sigmoid(self.fc2(out))
```

```
return out, attention_weights
```

*# Usage example with attention visualization*

```
def train_with_attention_visualization():
```

```
    model = SentimentLSTM(vocab_size=10000, embedding_dim=128,  
                           hidden_dim=64, num_layers=2)
```

*# Sample input*

```
sample_input = torch.randint
```